

Handwriting Recognition

Deep Learning Final Project

Amee Tan & Evan Chen

Agenda

1.

Recap of Project

2.

Results

3.

**Failures /
Methods Attempted**


4.

Successes

5.

Future Directions

Project Recap




331,059 training images

41,382 validation images



Handwriting Recognition using CRNN in Keras

Python notebook using data from [Handwriting Recognition](#) · 10,602 views · 1y ago · 🖨️ gpu



Pytorch

Model



Scoring Metrics

Correct Characters Predicted

On average, how many characters did we get correct in each name?

Prediction	A	I	E	K
Actual	A	L	E	X

50%

Correct Words Predicted

How many of them did we correctly predict?

Actual	Prediction
JOSE	JOSB
ANNIE	ANNIE
MARCOS	NAKOOS
MARIANNA	MARLANNA

1 out of 4

How did we do?

Correct characters
predicted

20%

Correct words
predicted

5

out of 41,382



Correct characters predicted : 82.16%

Correct words predicted : 69.10%

= 28,553 correct words

Failures / Methods Attempted

- Could not successfully load data into Google Collab without timing out
 - Trained in Kaggle instead
- Kaggle GPUs timed out when running overnight



Failures / Methods Attempted

- Could not successfully load data into Google Collab without timing out
 - Trained in Kaggle instead
- Kaggle GPUs timed out when running overnight

```
Epoch: 14
Train:
CTC Loss: 13.4237
Percent correct characters per word: 0.1979
Number of correct words: 5
Valid
CTC Loss 13.0708
Percent correct characters per word 0.2007
Number of correct words 1
```

```
Epoch 14/60
235/235 [=====]
loss: 5.6524 - val_loss: 5.2620
```


Failures/Methods Attempted

- Negative CTC loss values
 - Take LogSoftMax of predictions before feeding into CTC loss
- NaNs for loss when training on entire dataset

Epoch: 0

/opt/conda/lib/python3.7/site-packages/ipykernel: Change the call to include dim=X as an

Train:

CTC Loss: nan

Percent correct characters per word: 0.082

Number of correct words: 10

Valid

CTC Loss nan

Percent correct characters per word 0.1246

Number of correct words 5

```
train_loss = one_pass(model, dl_train, optimizer)
train_loss
```

```
tensor(-5.0396, grad_fn=<MeanBackward0>)
tensor(-4.1961, grad_fn=<MeanBackward0>)
tensor(6.1593, grad_fn=<MeanBackward0>)
tensor(22.7922, grad_fn=<MeanBackward0>)
tensor(-6.9500, grad_fn=<MeanBackward0>)
tensor(18.8461, grad_fn=<MeanBackward0>)
tensor(21.2995, grad_fn=<MeanBackward0>)
tensor(11.2788, grad_fn=<MeanBackward0>)
tensor(-6.8725, grad_fn=<MeanBackward0>)
tensor(10.3967, grad_fn=<MeanBackward0>)
tensor(20.5246, grad_fn=<MeanBackward0>)
tensor(18.8454, grad_fn=<MeanBackward0>)
tensor(14.8141, grad_fn=<MeanBackward0>)
tensor(6.6474, grad_fn=<MeanBackward0>)
tensor(-1.7920, grad_fn=<MeanBackward0>)
tensor(1.4204, grad_fn=<MeanBackward0>)
Avg loss 8.01088497042656
```

Learning Successes

Batch Size Experiments

	32	64
CTC Loss	24.9271	30.5298
Correct Characters Predicted	6.91%	6.36%
Correct Words Predicted	0	0

Learning Successes

Learning Rate Experiments

	0.01	0.001	Cosine Annealing Scheduler
CTC Loss	37.3579	24.5227	30.3909
Correct Characters Predicted	6.15%	6.91%	7.43%
Correct Words Predicted	0	0	1

Learning Successes

- MiSH Activation function
 - Longer to train vs. ReLU due to derivative calculation when $x < 0$
 - Better results

	ReLU	MiSH
Correct Characters Predicted	8.8%	20%
Correct Words Predicted	0	5

```
class Mish(nn.Module):
    def __init__(self):
        super().__init__()
    def forward(self, x):
        return (x * torch.tanh(F.softplus(x)))
```

Keras → Pytorch Translation Guide

Keras	Pytorch
(batch_size, img_length, img_width, channels)	(batch_size, channels, img_length, img_width)
<pre>inner = Conv2D(32, (3, 3), padding='same', name='conv1', kernel_initializer='he_normal')(input_data)</pre>	<pre>self.conv1 = nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3, padding=1)</pre>
<pre>inner = Dense(num_of_characters, kernel_initializer='he_normal', name='dense2')(inner)</pre>	<pre>self.linear2 = nn.Linear(in_features=1024, out_features=30)</pre>
<pre>inner = Bidirectional(LSTM(256, return_sequences=True), name = 'lstm1')(inner) inner = Bidirectional(LSTM(256, return_sequences=True), name = 'lstm2')(inner)</pre>	<pre>self.lstm1 = nn.LSTM(input_size=64, hidden_size=512, batch_first=True, bidirectional=True, num_layers=2)</pre>

Future Directions

- Train for more than 15-20 epochs on USF's GPUs
- Revise **percent correct** metric to be more forgiving
- Deeper dive into what was wrong with our metrics and loss function
- Try using a pre-trained model

 NVIDIA DEVELOPER



Text Recognition

Recognizes text from an image.

[VIEW MODELS >](#)

THANKS

CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, and infographics & images by **Freepik**